

高级数据结构

Advanced Data Structure

3. Access Methods for Multidimensional Data Retrieval 1

Review - Primary key access methods

■ Data stored on disk(s)


- Transfer unit: block (page)
- I/O complexity: in number of blocks accessed

■ Problem

- Minimize the number of transfer unit
- Handle dynamic database

■ Indexing

- Loading factor
- External memory based
- Dynamic database
- Easy implementation/maintenance



Review - B-tree family vs. Hashing

- *Hash-based* indices are best for exact match queries. Faster than *B+-tree*!
 - *Hash-based* indices typically require 1-2 I/Os per query
 - *B+-tree* requires 4-5 I/Os (logarithmic)
 - *B family tree* support answering
 - *range queries*
 - *nearest neighbor queries*
 - *ordered sequential scanning...*
- Hash-based* indices don't support.

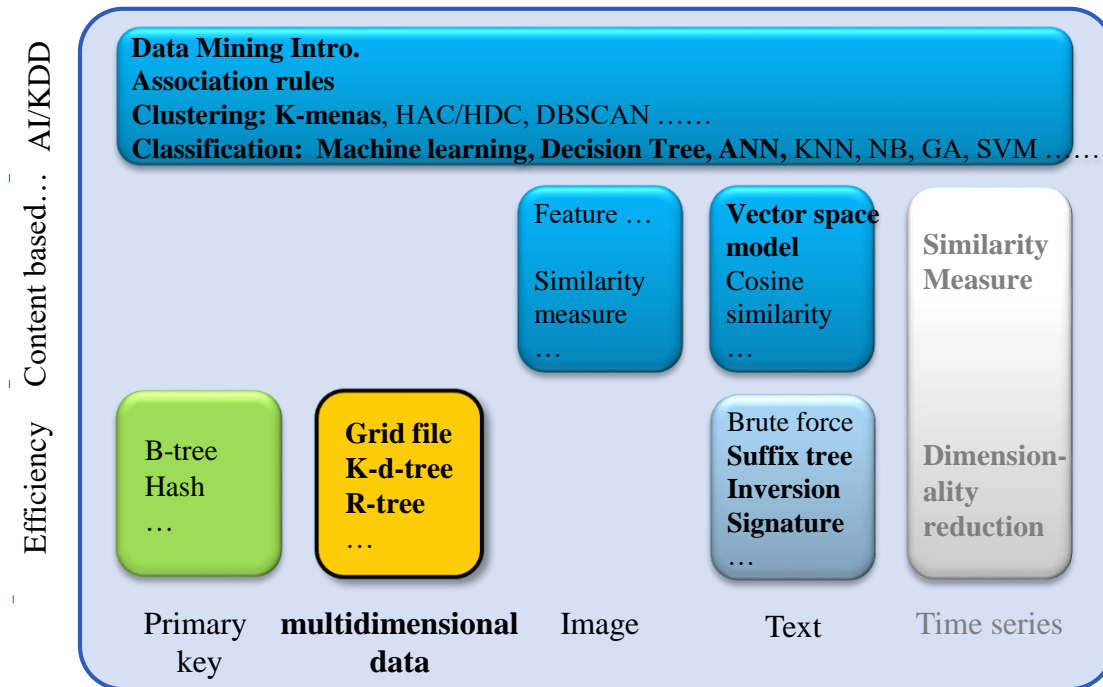
Agenda

■ Problem: **Multidimensional Data Retrieval**

■ Method:

- **Grid file**
- **K-d-tree**
- **R-tree**

.....



Agenda

- Problem: **Multidimensional Data Retrieval**
 - Point access
 - Spatial access
 - Key idea: Progressive refinement

- Method:
 - Partition space: Grid file, K-d-tree
 - Partition dataset: R-tree
 -
- Dimensionality Curse


Agenda

- Problem: **Multidimensional Data Retrieval**
- Method:
 - **Grid file**
 - **K-d-tree**
 - R-tree
 -

Motivating problem 1 - Spatial data accessing


- Given a point (region) set and a query, find the points (regions) that satisfy the query

Each building \rightarrow a point/polygon/...



Motivating problem 2 - Multi-key based accessing

Age > 18 and Age < 35

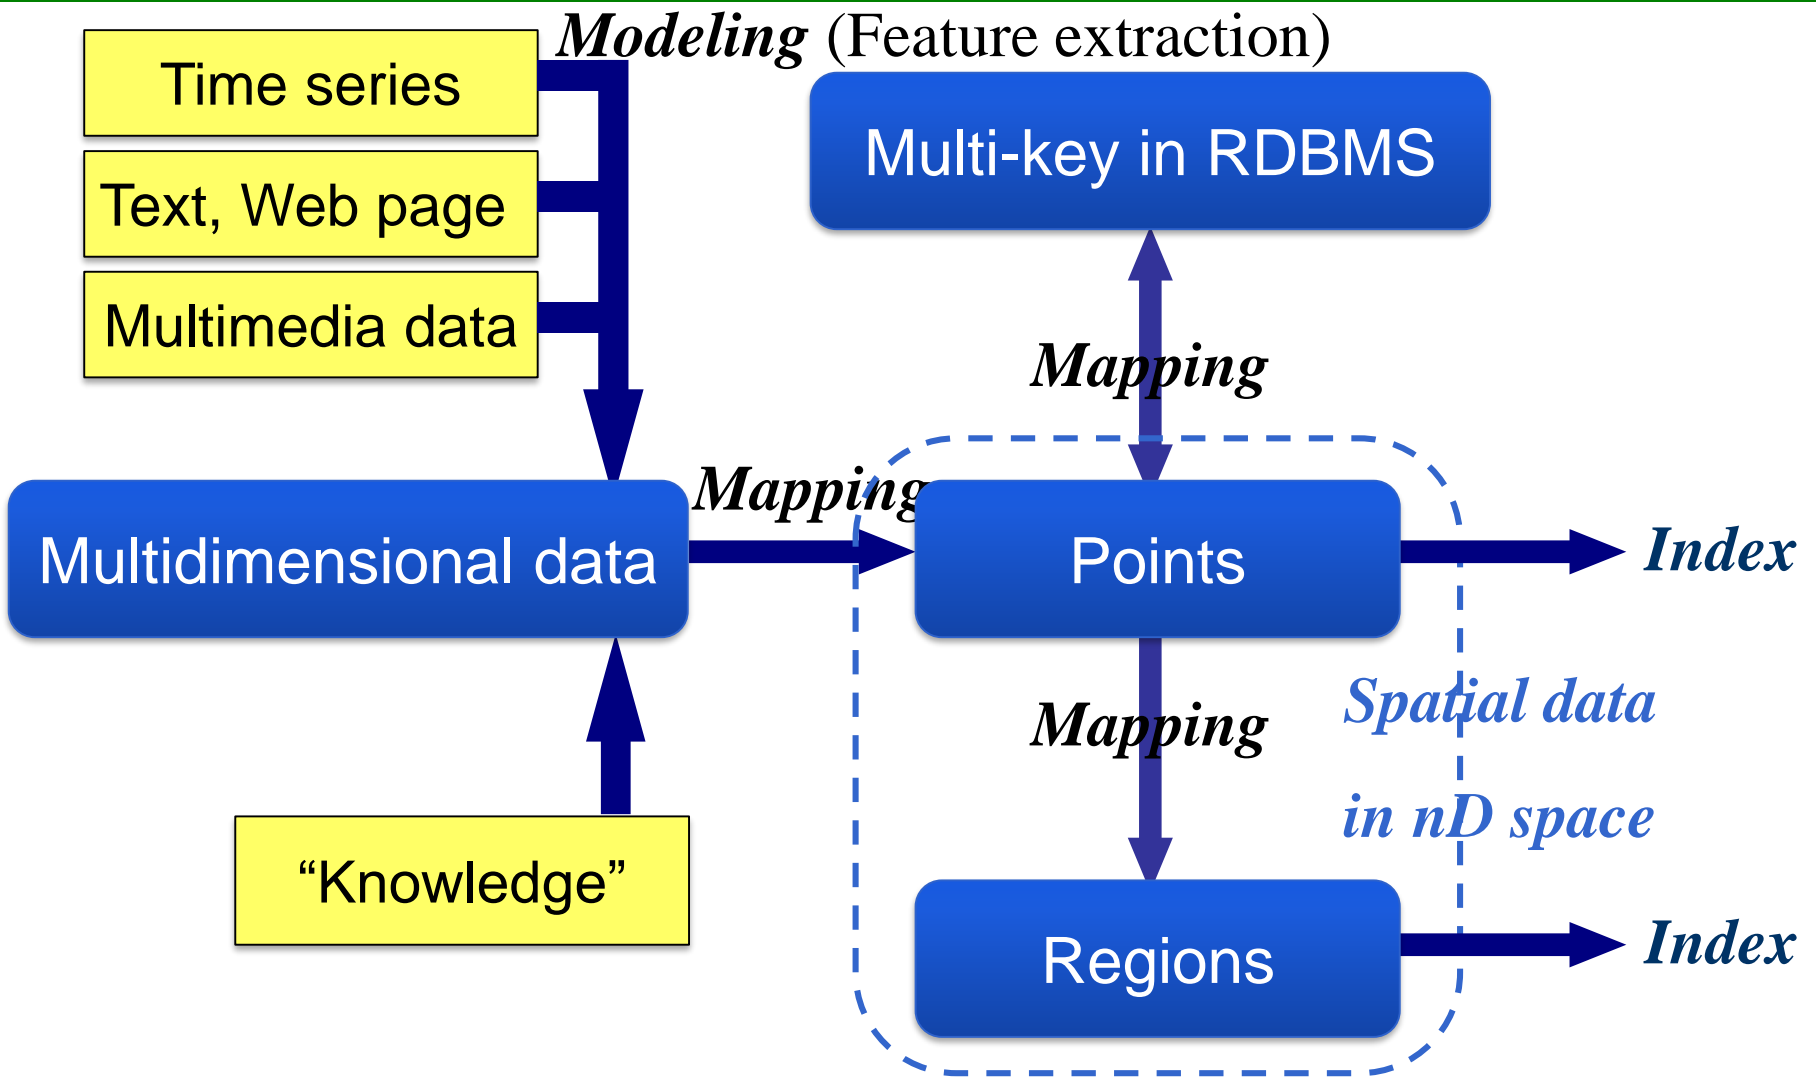


Map to 2D points (age,salary)

Age > 18 and Age < 35 and Salary > 10000

ID	Name	Age	Salary
123	S. John	33	30000
456	J. Tom	22	3000
...

Related topics – Searching multimedia data by content



Queries

Given

- data set $P=\{P_1, P_2, \dots\}$, $P_n=(x_n, y_n)$, or $P_n=(x_n, y_n, z_n)$,
- Distance measurement: $D(A, B)$

■ ***Exact match***

- Query point: $Q=(x, y)$
- $M(Q)=\{B \in P \mid B=Q\}$

■ ***Nearest neighbor queries***

- Query point: $Q=(x, y)$
- $NN(Q)=\{B \in P \mid \forall A \in P \rightarrow D(B, Q) \leq D(A, Q)\}$
- ***K-Nearest-neighbor queries***: find the k-nearest points to Q

■ ***Range queries*** (circle region)

- Query point: $Q=(x, y)$, threshold ε
- $RQ_C(Q)=\{B \in P \mid D(B, Q) \leq \varepsilon\}$


Range queries (rectangle region, ***window query***)

- Query region: $R=(x_1, y_1, x_2, y_2)$
- $RQ_R(Q)=\{B \in P \mid B \in R\}$


Distance measurement

$$L_p = \left(\sum_{i=1}^n |A_i - B_i|^p \right)^{1/p}$$

- L_2 : Euclidean distance
- L_1 : Manhattan, city-block distance
- $L_\infty(A, B) = \max_i |A_i - B_i|$



Methodologies



- Point Access Methods (PAMs): index point data
- Spatial Access Methods (SAMs): index both points and regions

Methodologies

- Idea: divide and conquer
 - Partition the search space into subspaces, how?
- Organize/partition the **embedding space**
 - →Grid-like structure for indexing spatial data?
- Organize/partition the **specific set of data**
 - →A tree-like structure for indexing spatial data?

Methods

- Inverted files
- Grid files
- K-d-trees
- Space filling curves
- R-trees


Agenda

- Problem: **Multidimensional Data Retrieval**
- Method:
 - **Grid file:** Generalization of extendible hashing in multiple dimensions
 - **K-d-tree**
 - R-tree
 -

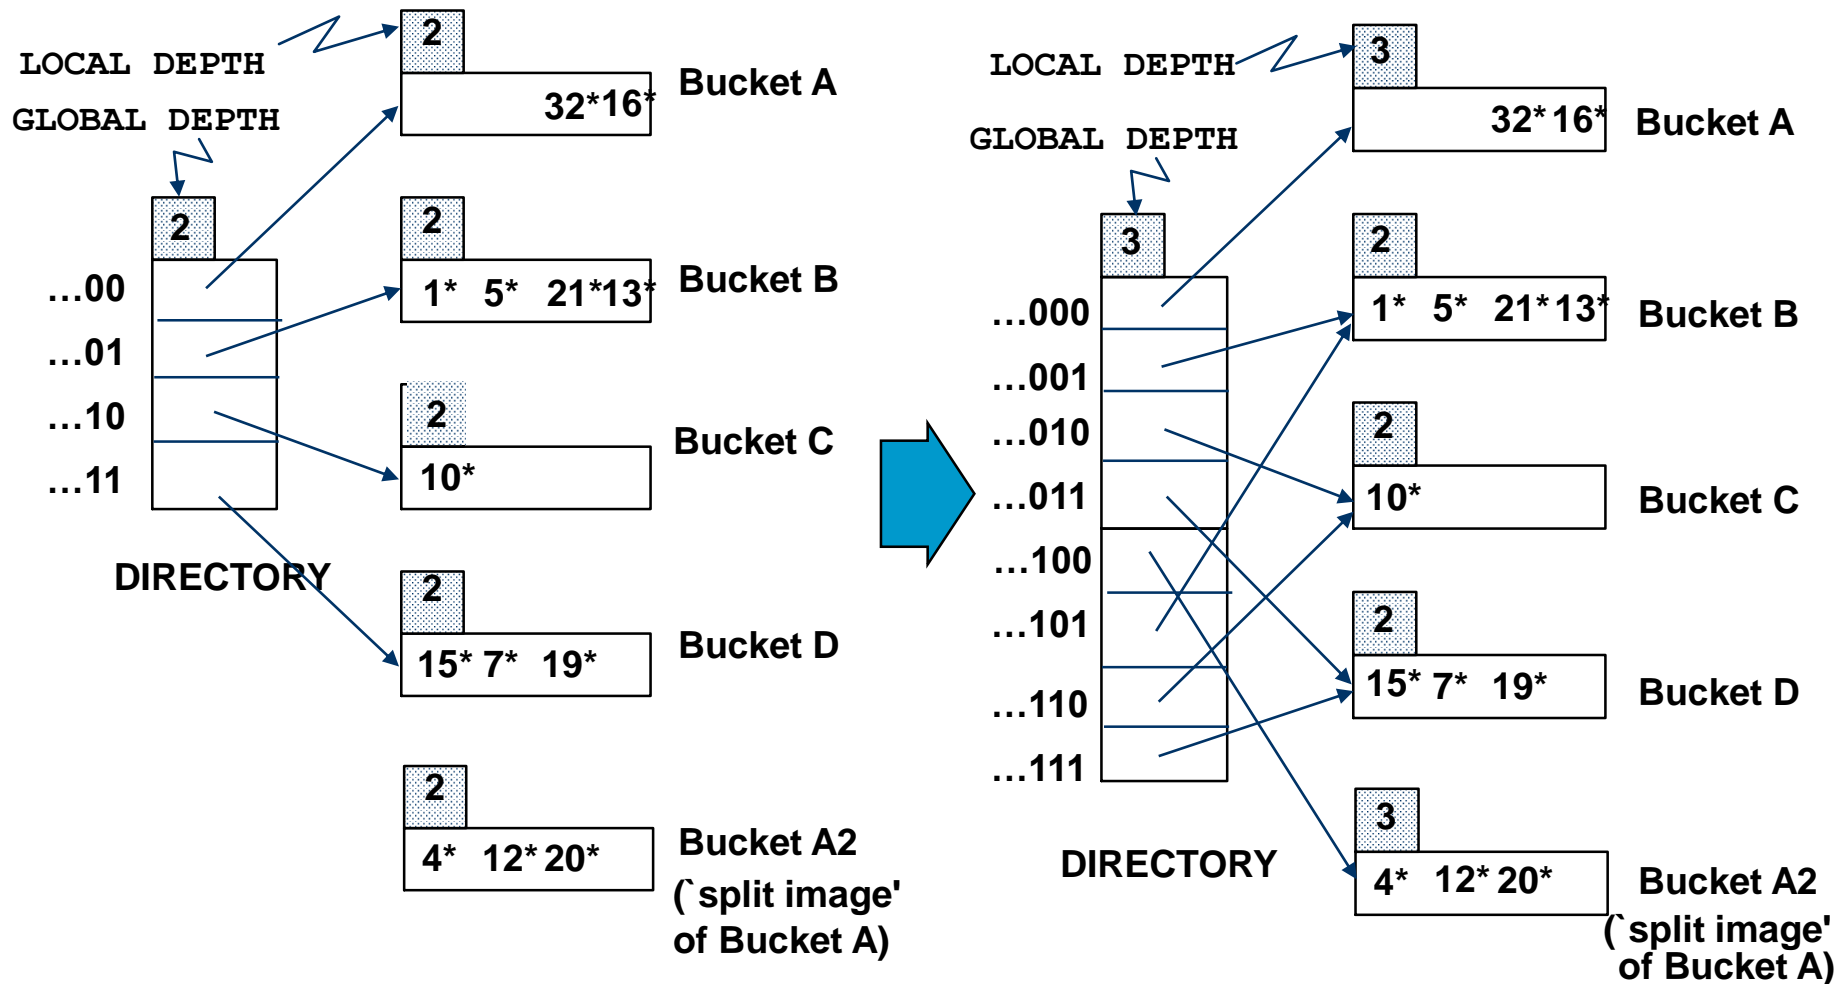
Grid - Naïve version

- Idea: Use a grid to partition the space into cells
- All cell with fixed size
- Non-uniform distribution → **Uncertain searching depth**


Ideally, we want each cell associated with one page





Extendible hashing



Grid File



Grid File - Example




Grid file structure

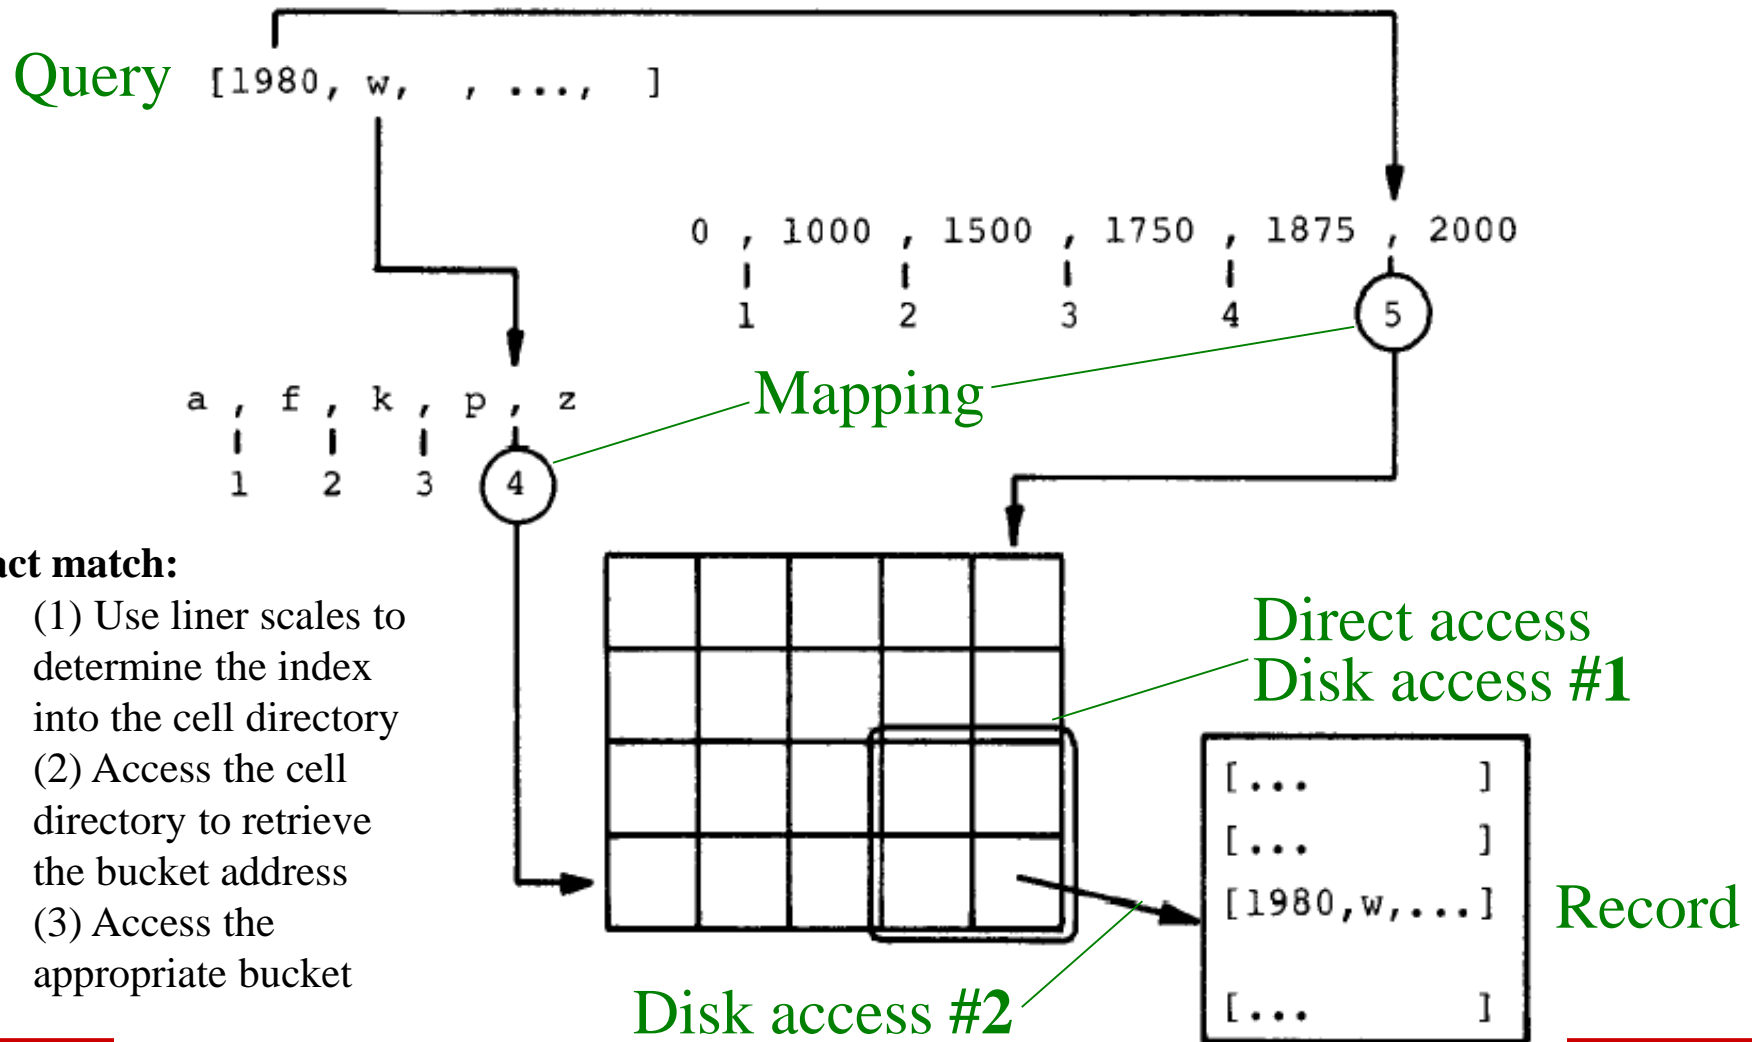
- Data stored in *disk blocks (buckets)*
- Dynamic structure using a grid directory
 - **Grid array (directory)**: a 2 dimensional array with pointers to buckets
 - $G(0 \dots n_x-1, 0 \dots n_y-1)$
 - stored in *disk blocks or main memory*
 - **Linear scales**: Two 1-dimensional arrays that used to access the grid array
 - $X(0 \dots n_x-1), Y(0 \dots n_y-1)$
 - E.g. $X=(1600,1900,1950,1980,1990,2010)$
 - stored in *main memory*

Operations on grid file

- Mapping
 - $(x,y) \rightarrow (cx,cy)$
- Direct access
 - $G(cx,cy) = \text{pointer to bucket}$
- Adjust directory
 - Merge
 - Split




Operations – Mapping & direct access (exact match)



Grid file search - Range queries


■ Range queries:

- (1) Use linear scales to determine the index into the cell directory
- (2) Access the cell directory to retrieve the bucket addresses to visit
- (3) Access the buckets
- (4) Verification



Grid file search – Nearest neighbor query - I


- Identify the bucket to which the query point Q belongs.
- By examining the corresponding records, generate:
(1) a *temporary NN point* R' and (2) a *temporary NN distance* $r' = D(Q, R')$.



Grid file search – Nearest neighbor query - II


■ Select candidate

- Compute the minimum distances $\text{MinDist}(Q, C_i)$ from the query point Q to each cell C_i intersecting with the *temporary NN sphere*
- For each cell C_i , if $\text{MinDist}(Q, C_i) \geq r$, it is discarded (why?)
- Otherwise, the cell is added to a *priority candidate list*, in which the cell with smaller *MinDist* has a higher priority



Grid file search – Nearest neighbor query - III





- The disk block of each candidate cell is accessed to update both *temporary NN point* R' and NN distance r' .
- The search stops whenever *MinDist* of the next candidate \geq the current NN distance r' .




Grid file insertion

- Determine the bucket into which insertion will occur.
- If space in bucket, insert.
- Else, split bucket.
 - If it cause the directory split, do so and adjust linear scales.
- Directory need split?
- Which dimension to split?
- Insertion of new entries potentially requires a **complete reorganization** of the cell directory

Example – 2D



Example – 3D




Grid File deletion

- Determine the bucket from which deletion will occur.
- Delete the record from the bucket
- If *space utilization* of the bucket is less than a threshold
 - Merge the bucket with adjacent one

Remarks



- A **space partitioning** strategy
- **Select** dividers (**not fixed**) along each dimension.
- Adapt to **non-uniform distributions**
- **Two disk access** (exact match)
- What if the attribute values are **correlated** (e.g., age and salary)
- **Cell directory potentially exponential in the number of dimensions**
- **Split and merge may cause a rewrite of the entire directory**



Agenda

- Problem: **Multidimensional Data Retrieval**
- Method:
 - **Grid file**
 - **K-d-tree**: Generalization of binary search tree in multiple dimensions
 - R-tree
 -

Binary search tree



15, 7, 30,.....



K-d tree: Extension of Binary Search Tree

K-d tree structure

- At each Internal node, the k-d tree divides the k-dimensional space into two parts by a (k-1)-dimensional *divider*

height 0 cuts on x, height 1 cuts on y,



height 2 cuts on x, height 3 cuts on y, ...



x coordinate first

K-d tree structure



- Internal nodes N : k -dimensional data
 - N is a decision point on a specific dimension
 - Splits the $1 + (\text{level}(N) \bmod k)$ dimension into two regions
 - 2d-node $\{x, y: \text{real}, \text{info}: \dots, \text{left}, \text{right}: \wedge 2\text{d-node}\}$
- Leaf nodes: failure



- Data stored in main-memory



K-d tree search - Exact match

- Straightforward!





K-d tree search – range queries

- Straightforward!





K-d tree insert

- Straightforward!





K-d tree deletion



- Deletion may cause re-organization of the tree under the deleted node
- Deletion of a node n_1 :
 - Locate the node n_1 on the tree
 - Find a *candidate* replacement node n_2
 - Replace the contents of n_1 with n_2
 - Recursively delete n_2
- Candidate node:
 - If n_1 is an x node, then a possible candidate is
 - a node with the smallest x value in the right subtree, or
 - a node with the largest x value in the left subtree
 - Similarly for y nodes

K-d tree deletion



Deletion:



Locate the node n1 on the tree

Find a candidate replacement node n2

Replace the contents of n1 with n2

Recursively delete n2

K-d tree deletion



Deletion:



Locate the node $n1$ on the tree

Find a candidate replacement node $n2$

Replace the contents of $n1$ with $n2$

Recursively delete $n2$

K-d tree deletion





Deletion:

- Locate the node n1 on the tree**
- Find a candidate replacement node n2**
- Replace the contents of n1 with n2**
- Recursively delete n2**

Remarks – K-d tree



- Extension of Binary Search Tree
- Height of the tree: $O(\log n)$
Search time for exact match: $O(\log n)$
- **K-d tree**
 - Depends heavily on the insertion order of the points (best/worst case?)
 - The divider does not divide the space at the best possible positions, resulting in an unbalanced tree
 - Uncertain searching depth
- The tree scales linearly with the number of dimensions
- Main memory based
- Elegant ideas that
 - have been widely used in various fields
 - appear in many text books
- Many extensions
 - Adaptive k-d tree
 - K-d-B tree



Adaptive k-d tree

- An improved version of k-d tree
 - **All points are stored in the leaves**
 - Internal nodes contain the dimension (e.g. x or y), and the coordinate of the respective split
 - The divider is chosen as the **median value** between the points
 - The partitioning dimension is chosen **adaptively** (not alternately)
- It works well when the data is known a-priori, and there are rare updates in the tree



Adaptive k-d tree



k-d-B tree

- The k-d-B tree combines properties of both the adaptive k-d tree and the B-tree
 - More than one divider divide a tree node (depending on the tree's storage utilization)
 - All nodes of the tree correspond to disk blocks
 - Like the B-tree, the k-d-B tree is perfectly balanced
 - Insertion becomes quite complex and expensive
 - Storage utilization guarantee?

k-d-B tree



Conclusions

- Problem: Accessing method for multidimensional data
 - Data: spatial points/ regions, RDBMS,
 - Several query forms
- Grid file and k-d-tree
 - Architecture, update, access ...
 - Grid file and k-d-tree are used in many stand-alone systems

Homework

- Read 4.1 (optional), 4.2.1, 4.2.2 and additional reading material
- Given a 2D data set by yourself, construct the Grid file and k-d tree

Thanks

Feedback welcome